



MS&E 228: Applied Causal Inference Powered by ML and AI

Lecture 8: Deep Dive into OLS Coefficients and the Double Machine Learning Estimator

Vasilis Syrgkanis

Stanford University

Winter 2026

Readings: *Applied Causal Inference Powered by ML and AI*, §1, 4, 9.

Goals for Today

1. **Develop deeper intuition for OLS coefficients:** What does a coefficient in a multivariate regression really mean? When is it uniquely identified, and what drives its variance?
2. **Discover another “insensitive” formula:** Can we simplify inference for OLS coefficients by ignoring nuisance estimation, similar to the DR estimator?
3. **Move beyond full linearity:** How can we allow flexible ML models for covariates while still getting valid causal inference on treatment effects?
4. **Build a practical algorithm for continuous treatments:** What estimator can we use in practice, and what does it estimate when treatment effects are heterogeneous?

Part I: Understanding OLS Coefficients

The Frisch-Waugh-Lovell Theorem

Where We Left Off

Last lecture, under full linearity $\mathbb{E}[Y(d)|X] = \theta d + \beta'X$, all causal estimands reduce to estimating θ .

We posed four key questions:

1. How can we interpret θ intuitively?
2. When is θ uniquely identified in the best linear predictor?
3. What drives high variance in $\hat{\theta}$?
4. What is the analogue of “overlap” under linearity?

Today we answer all four questions through two marvelous theorems!

The Frisch-Waugh-Lovell Theorem

Theorem (Frisch-Waugh-Lovell)

The coefficient $\hat{\theta}$ on D in $\text{OLS}(Y \sim D, X)$ is identical to the coefficient from:

Step 1: Regress Y on X and compute residuals—the “surprise” in Y :

$$\check{Y} = Y - \text{OLS}(Y \sim X).\text{predict}(X)$$

Step 2: Regress D on X and compute residuals—the “surprise” in D :

$$\check{D} = D - \text{OLS}(D \sim X).\text{predict}(X)$$

Step 3: Run simple line-fit on the surprises:

$$\hat{\theta} = \text{OLS}(\check{Y} \sim \check{D}) \quad \Leftrightarrow \quad \text{solve empirical normal equations } \mathbb{E}_n[(\check{Y} - \hat{\theta}\check{D})\check{D}] = 0$$


Visual Intuition for FWL

Original Problem

OLS($Y \sim D, X$)
High-dimensional
regression

Same $\hat{\theta}$

FWL Equivalent



A complex multivariate regression reduces to a simple line fit on residuals!

The FWL Formula

From the simple line-fit OLS($\check{Y} \sim \check{D}$), we get:

$$\hat{\theta} = \frac{\mathbb{E}_n[\check{Y} \cdot \check{D}]}{\mathbb{E}_n[\check{D}^2]} = \frac{\text{Cov}_n(\check{Y}, \check{D})}{\text{Var}_n(\check{D})}$$

Key Insight

$\hat{\theta}$ is the **normalized correlation** between residualized outcome and residualized treatment—how much the “surprise” in Y moves with the “surprise” in D .

Think of it as: When treatment is higher than expected given confounders, is the outcome also higher than expected?

Python Pseudocode: The FWL Theorem

```
# --- Standard OLS approach ---
model_full = LinearRegression().fit(np.c_[D, X])
theta_full = model_full.coef_[0]

# --- FWL three-step procedure ---
# 1. Residualize Y: Regress Y on X
Yres = Y - LinearRegression().fit(X, Y).predict(X)
# 2. Residualize D: Regress D on X
Dres = D - LinearRegression().fit(X, D).predict(X)
# 3. Final Regression: Regress residualized Y on residualized D
theta_fwl = LinearRegression(fit_intercept=False).fit(Dres, Yres)

# --- Verification ---
print(f"Full OLS theta: {theta_full:.6f}")
print(f"FWL OLS theta: {theta_fwl:.6f}")
```

Poll: Interpreting OLS Coefficients

Poll Question

What does the coefficient θ on the first feature in $\text{OLS}(Y \sim X)$ represent?

- A. The first coordinate of $\mathbb{E}[XX']^{-1}\mathbb{E}[XY]$
- B. The correlation between Y and the first feature
- C. The normalized correlation between the residual of Y and the residual of the first feature, after removing linear predictions from other variables
- D. The average effect of changing the first feature by one unit



(Poll Everywhere)

Live Coding: Verifying the FWL Theorem


Exercise

Goal: Verify that $\text{OLS}(Y \sim D, X)$ gives the same coefficient as the FWL process.

Steps:

1. Run full OLS: $Y \sim D + X$
2. Regress Y on X , get residuals $\check{Y}$
3. Regress D on X , get residuals $\check{D}$
4. Regress $\check{Y}$ on $\check{D}$
5. Compare coefficients!

Dataset: Price elasticity simulation



Open notebook and follow along!

The Frisch-Waugh-Lovell Theorem (in Population Limit)

Theorem (Frisch-Waugh-Lovell) in population

The coefficient θ on D in the Best Linear Predictor $\text{BLP}(Y \sim D, X)$ is identical to the coefficient from:

Step 1: Calculate BLP of Y using X and compute *population* residuals:

$$\tilde{Y} = Y - \text{BLP}(Y \sim X).\text{predict}(X)$$

Step 2: Calculate BLP of D using X and compute *population* residuals:

$$\tilde{D} = D - \text{BLP}(D \sim X).\text{predict}(X)$$

Step 3: Find coefficient in BLP of $\tilde{Y}$ using $\tilde{D}$:

$$\theta = \text{coeff in BLP}(\tilde{Y} \sim \tilde{D}) \Leftrightarrow \text{solve normal equations } \mathbb{E}[(\tilde{Y} - \theta\tilde{D})\tilde{D}] = 0$$

Answering Question 1: Interpreting θ

Question: How can we better interpret the population coefficient θ that OLS converges to?

$$\theta = \frac{\mathbb{E}[\tilde{Y}\tilde{D}]}{\mathbb{E}[\tilde{D}^2]} = \frac{\text{Cov}(\tilde{Y}, \tilde{D})}{\text{Var}(\tilde{D})}$$

Answer

It is the slope of the line-fit between:

- ▶ The part of Y not linearly predictable from X (the “surprise” in outcome)
- ▶ The part of D not linearly predictable from X (the “surprise” in treatment)

θ captures how much the (*linearly*) unexplained variation in Y moves with the (*linearly*) unexplained variation in D , after controlling for X .

Much more intuitive than “the first coordinate of $\mathbb{E}[XX']^{-1}\mathbb{E}[XY]$ ”!

Answering Question 1: When is θ Uniquely Identified?

Question: When does the population BLP have a unique solution for θ ?

Answer

When D cannot be perfectly linearly predicted from X —i.e., when there's some “surprise” left in treatment.

- ▶ If $D = \gamma'X$ exactly (perfect collinearity), then $\tilde{D} = 0$ and $\text{Var}(\tilde{D}) = 0$
- ▶ The coefficient θ is undefined—this is the **collinearity problem**
- ▶ But this is much more benign than the overlap assumption!
- ▶ We only need $\text{Var}(\check{D}) > 0$ **on average**, not pointwise overlap

Part II: Proof of the FWL Theorem

Understanding Why It Works

OLS as Orthogonal Decomposition

Key Insight

Finding $\text{OLS}(Y \sim D, X)$ coefficients is equivalent to finding a decomposition:

$$Y = \theta D + \beta' X + c + \epsilon$$

where ϵ is empirically orthogonal to $(1, D, X)$ (denoted as $\epsilon \perp_n (1, D, X)$), i.e.,

$$\mathbb{E}_n[\epsilon] = 0, \quad \mathbb{E}_n[\epsilon D] = 0, \quad \mathbb{E}_n[\epsilon X] = 0$$

Any (coefficients, residual) pair satisfying this orthogonality is a valid $\text{OLS}(Y \sim D, X)$ solution. *Think of it as: OLS extracts everything useful from the predictors. The residual has no remaining correlation.*

Strategy: We will show that the FWL procedure produces coefficients satisfying these normal equations.

The Residualization Operator

Define the **residualization operator**: For any variable V ,

$$\breve{V} = V - \gamma'_V X - c_V$$

where γ_V, c_V are chosen so that $\breve{V} \perp_n (1, X)$.

This removes the part of V that's linearly predictable from X .

Key Property: Additivity

$$\widetilde{A + B} = \breve{A} + \breve{B}$$

Proof: Write $A + B = (\gamma_A + \gamma_B)'X + (c_A + c_B) + (\breve{A} + \breve{B})$.

Since $\breve{A} + \breve{B} \perp_n (1, X)$, this is a valid decomposition. (*The surprise in $A + B$ is the sum of the surprises.*)

Proving FWL

Start with OLS decomposition: $Y = \theta D + \beta' X + c + \epsilon$, with $\epsilon \perp_n (1, D, X)$

Residualize both sides:

$$\check{Y} = \theta \check{D} + \beta' \check{X} + \check{c} + \check{\epsilon}$$

Now observe:

- ▶ $\check{X} = 0$ (X is perfectly linearly predictable from X)
- ▶ $\check{c} = 0$ (constant is perfectly linearly predictable)
- ▶ $\check{\epsilon} = \epsilon$ ($\epsilon \perp_n (1, X)$ means ϵ is its own residual)

Therefore:

$$\check{Y} = \theta \check{D} + \epsilon, \quad \text{with } \epsilon \perp_n \check{D}$$

Thus, the pair θ, ϵ is a valid OLS($\check{Y} \sim \check{D}$) decomposition and hence the solution to the OLS($\check{Y} \sim \check{D}$) problem! □

Part III: Variance and Standard Errors

The Insensitivity Theorem

The Insensitivity Theorem

Theorem (Insensitivity)

When calculating standard errors for $\hat{\theta}$, we can ignore the estimation error in the first two OLS steps.

Asymptotically, $\hat{\theta}$ behaves as if we had the true population residuals:

$$\tilde{D} = D - \text{BLP}(D|X), \quad \tilde{Y} = Y - \text{BLP}(Y|X)$$

You might worry: “I estimated those residuals—doesn’t that add uncertainty?”
Surprisingly, no!

Implication: For inference, we only need to analyze a simple line-fit. Same “insensitivity” as the DR estimator!

Deriving the Asymptotic Distribution

Starting from the population residuals:

$$\begin{aligned}\sqrt{n}(\hat{\theta} - \theta) &= \sqrt{n} \left(\frac{\mathbb{E}_n[\check{D}\check{Y}]}{\mathbb{E}_n[\check{D}^2]} - \theta \right) \approx \sqrt{n} \left(\frac{\mathbb{E}_n[\tilde{D}\tilde{Y}]}{\mathbb{E}_n[\tilde{D}^2]} - \theta \right) = \sqrt{n} \frac{\mathbb{E}_n[\tilde{D}\tilde{Y}] - \theta \mathbb{E}_n[\tilde{D}^2]}{\mathbb{E}_n[\tilde{D}^2]} \\ &= \sqrt{n} \frac{\mathbb{E}_n[\tilde{D}(\tilde{Y} - \theta \tilde{D})]}{\mathbb{E}_n[\tilde{D}^2]}\end{aligned}$$

By LLN, denominator $\rightarrow \text{Var}(\tilde{D})$. By CLT:

$$\boxed{\sqrt{n}(\hat{\theta} - \theta) \xrightarrow{d} N \left(0, \frac{\mathbb{E}[\tilde{D}^2(\tilde{Y} - \theta \tilde{D})^2]}{\text{Var}(\tilde{D})^2} \right)}$$

Computing Standard Errors in Practice

The asymptotic variance can be estimated using empirical analogues:

$$\hat{\sigma}^2 = \frac{\mathbb{E}_n[\check{D}^2(\check{Y} - \hat{\theta}\check{D})^2]}{\text{Var}_n(\check{D})^2}$$

Standard error: $\text{s.e.} = \hat{\sigma}/\sqrt{n}$

Practical Note

This is exactly the **HC0 robust standard error** that OLS packages compute! They don't calculate it this way, but it's numerically equivalent.

Bottom line: just use robust standard errors from your favorite package—they're doing the right thing.

Python Pseudocode: Standard Errors via the Insensitivity Theorem

```
# Final stage residuals:
# outcome residual minus predicted from treatment residual
epsilon = Y_residuals - theta_fwl * D_residuals

# Asymptotic variance formula:
#  $\sigma^2 = E[(Y_{res} - \theta * D_{res})^2 * D_{res}^2] / E[D_{res}^2]^2$ 
numerator = np.mean(epsilon**2 * D_residuals**2)
denominator = np.mean(D_residuals**2)**2
sigma_sq = numerator / denominator

# Standard error = sigma / sqrt(n)
se_manual = np.sqrt(sigma_sq / len(Y))
```


Live Coding: Computing Standard Errors

Exercise

Goal: Verify that our manual SE formula matches statsmodels HC0.

Steps:

1. Get SE from statsmodels (HC0)
2. Compute residuals $\check{Y}, \check{D}$ from Part 1
3. Calculate $\epsilon = \check{Y} - \hat{\theta}\check{D}$
4. Apply the variance formula
5. Compare: $\text{s.e.} = \hat{\sigma}/\sqrt{n}$



Open notebook and follow along!

Answering Question 3: Drivers of High Variance

Question: What are the drivers of high variance for $\hat{\theta}$?

From $\hat{\theta} = \text{Cov}_n(\check{Y}, \check{D}) / \text{Var}_n(\check{D})$, we see two key drivers:

1. **Small denominator:** If $\text{Var}(\check{D})$ is small (treatment is highly predictable from confounders), we divide by a small number $\Rightarrow$ high variance
Little “surprise” in treatment means little to learn from.
2. **Large numerator noise:** If the residual $\check{Y} - \theta\check{D}$ has high variance, the estimate is noisy

Key Insight

High linear predictability of treatment from covariates leads to high variance—but this is much weaker than requiring overlap everywhere!

Answering Question 4: The Analogue of Overlap

Question: What is the analogue of “overlap” under linearity?

Answer

$$\text{Var}(\tilde{D}) > \epsilon \text{ for some } \epsilon > 0$$

- ▶ We need the residual treatment variance to be bounded away from zero
- ▶ This is an **average condition**, not a pointwise condition
- ▶ Even if D is perfectly predictable in some regions of X , we're fine as long as it's unpredictable in other regions

Compare to binary treatment overlap ($0 < P(D = 1|X) < 1$ everywhere)—linearity lets us extrapolate!

Poll 2: Variance Drivers

Poll Question

In $\text{OLS}(Y \sim D, X)$, when would you expect high variance for the coefficient on D ?

- A. When Y has high variance
- B. When D can be well predicted linearly from X
- C. When X has many dimensions
- D. When the sample size is small



(Poll Everywhere)

Part IV: Why Insensitivity Holds

The Same Logic as Doubly Robust

The Insensitivity Argument

Our target parameter solves: $\mathbb{E}[(\tilde{Y} - \theta\tilde{D})\tilde{D}] = 0$

Define for arbitrary first-stage coefficients γ_Y, γ_D :

$$M(\theta, \gamma_Y, \gamma_D) = \mathbb{E}[(Y - \gamma_Y'X - \theta(D - \gamma_D'X))(D - \gamma_D'X)]$$

Claim

This formula is **insensitive** to γ_Y and γ_D at the true values—small errors in the first stage don't affect the final answer.

Same logic as the DR estimator!

Checking Insensitivity

Partial derivative w.r.t. γ_Y :

$$\partial_{\gamma_Y} M(\theta_0, \gamma_Y^*, \gamma_D^*) = -\mathbb{E}[X(D - \gamma_D^{*'} X)] = -\mathbb{E}[X\tilde{D}] = 0$$

The last equality holds because $\tilde{D} \perp X$ by construction!

Similarly, $\partial_{\gamma_D} M = 0$ by the orthogonality properties.

Conclusion

First-stage estimation errors don't affect the asymptotic distribution of $\hat{\theta}$. We can use simple standard error formulas!

Summary: Four Questions Answered

Question	Answer
When is θ uniquely identified?	When D is not perfectly linearly predictable from X (some “surprise” remains)
How to interpret θ ?	Slope of line-fit between residualized Y and residualized D (correlation of surprises)
What drives high variance?	High noise in Y , high linear predictability of D from X (little surprise)
Analogue of overlap?	$\text{Var}(\tilde{D}) > \epsilon > 0$ (average condition, not pointwise)

These answers come from the FWL theorem + Insensitivity theorem!

Part V: Beyond Full Linearity

Partial Linearity and Double Machine Learning

The Partial Linear Model

Full linearity is restrictive. What if we relax it?

Partial Linear Model

$$Y = \theta D + f(X) + \epsilon, \quad \mathbb{E}[\epsilon|D, X] = 0$$

- ▶ $f(X)$ can be **any function**—no linearity assumed for how confounders affect outcome
- ▶ Only the treatment effect is assumed constant: $\mathbb{E}[Y(d) - Y(0)|X] = \theta d$
- ▶ This is much more flexible while still enabling valid inference on θ

We're saying: "I don't know how confounders affect the outcome—could be complicated. But each unit of treatment has the same effect."

The Generalized FWL Idea

What if we replace “linearly predictable” with “predictable” in FWL?

Generalized Double Residualization:

1. Learn $\hat{h}(X)$ predicting Y from X using **any ML method**
2. Learn $\hat{p}(X)$ predicting D from X using **any ML method**
3. Compute residuals (surprises): $\check{Y} = Y - \hat{h}(X)$, $\check{D} = D - \hat{p}(X)$
4. Run simple line-fit: $\hat{\theta} = \text{OLS}(\check{Y} \sim \check{D})$

Question: Does this recover the correct θ ?

Population Identification

In the population, with true conditional expectations $h_0(X) = \mathbb{E}[Y|X]$ and $p_0(X) = \mathbb{E}[D|X]$:

$$\begin{aligned}\tilde{Y} &= Y - \mathbb{E}[Y|X] = \theta D + f(X) + \epsilon - \theta \mathbb{E}[D|X] - f(X) \\ &= \theta(D - \mathbb{E}[D|X]) + \epsilon = \theta \tilde{D} + \epsilon\end{aligned}$$

Since $\mathbb{E}[\epsilon|D, X] = 0$, we have $\mathbb{E}[\epsilon|\tilde{D}] = 0$.

The $f(X)$ terms cancel out! We don't need to know what f is.

Conclusion

θ is the solution to OLS($\tilde{Y} \sim \tilde{D}$) in the population!

Does Insensitivity Still Hold?

Define for arbitrary predictive models h, p :

$$M(\theta, h, p) = \mathbb{E}[(Y - h(X) - \theta(D - p(X)))(D - p(X))]$$

Check sensitivity to h : Let $h_t = h_0 + t\Delta$

$$\partial_t M(\theta, h_t, p_0)|_{t=0} = -\mathbb{E}[\Delta(X)(D - p_0(X))] = -\mathbb{E}[\Delta(X)\mathbb{E}[D - p_0(X)|X]] = 0$$

Check sensitivity to p : Similarly equals 0.

Yes!

The formula is insensitive to the predictive models—even complex ML models!
We can use powerful ML for confounders and still trust our standard errors.

The Overfitting Problem

Wait!

There's a catch we saw with DR: **overfitting and double-dipping!**

If we use the same data to:

1. Train $\hat{h}$ and $\hat{p}$
2. Compute residuals
3. Fit the final regression

We may overfit and get biased estimates. (*The model “memorizes” training points, making residuals artificially small.*)

Solution

Cross-fitting, just like with DR! For each point, predict using a model trained on *other* points.

The Double Machine Learning Algorithm

Algorithm (Double ML with K-fold Cross-Fitting)

1. Split data into K folds
2. For each fold k :
 - ▶ Train $\hat{h}^{(-k)}(X)$ on out-of-fold data to predict Y from X
 - ▶ Train $\hat{p}^{(-k)}(X)$ on out-of-fold data to predict D from X
 - ▶ For each sample i in fold k : $\check{Y}_i = Y_i - \hat{h}^{(-k)}(X_i)$, $\check{D}_i = D_i - \hat{p}^{(-k)}(X_i)$
3. Run simple OLS on all residuals (no intercept):

$$\hat{\theta} = \frac{\mathbb{E}_n[\check{Y}\check{D}]}{\mathbb{E}_n[\check{D}^2]}$$

Python Pseudocode for Double ML

```
from sklearn.model_selection import KFold

def double_ml(Y, D, X, n_folds=5):
    kf = KFold(n_splits=n_folds, shuffle=True)
    Y_res, D_res = np.zeros(len(Y)), np.zeros(len(Y))

    for train_idx, test_idx in kf.split(X):
        # Train on out-of-fold, predict for in-fold;
        # AutoML is placeholder for any ML method
        h_model = AutoML().fit(X[train_idx], Y[train_idx])
        Y_res[test_idx] = Y[test_idx] - h_model.predict(X[test_idx])

        p_model = AutoML().fit(X[train_idx], D[train_idx])
        D_res[test_idx] = D[test_idx] - p_model.predict(X[test_idx])

    # Final line-fit on surprises
    theta_hat = np.mean(Y_res * D_res) / np.mean(D_res**2)
    return theta_hat
```

Asymptotic Normality of Double ML

Theorem

Under regularity conditions, the cross-fitted DML estimator is asymptotically normal:

$$\sqrt{n}(\hat{\theta} - \theta_0) \xrightarrow{d} N\left(0, \frac{\mathbb{E}[(\tilde{Y} - \theta \tilde{D})^2 \tilde{D}^2]}{\text{Var}(\tilde{D})^2}\right)$$

Required conditions (the ML models need to be “good enough”):

1. Product rate: $\sqrt{n} \cdot \text{RMSE}(\hat{h}) \cdot \text{RMSE}(\hat{p}) \rightarrow 0$
2. Treatment model rate: $n^{1/4} \cdot \text{RMSE}(\hat{p}) \rightarrow 0$
3. Consistency: $\text{RMSE}(\hat{h}) \rightarrow 0$

We saw already how Random forests, Lasso, Neural Nets typically satisfy these conditions under small “effective dimension” assumptions.

Standard Errors for Double ML

Estimate the asymptotic variance using empirical analogues:

$$\hat{\sigma}^2 = \frac{\mathbb{E}_n[(\check{Y} - \hat{\theta}\check{D})^2\check{D}^2]}{\mathbb{E}_n[\check{D}^2]^2}$$

Standard error: $\text{s.e.} = \hat{\sigma}/\sqrt{n}$

```
def double_ml_with_se(Y, D, X, n_folds=5):
    theta_hat, Y_res, D_res = double_ml(Y, D, X, n_folds)

    # Compute standard error
    epsilon = Y_res - theta_hat * D_res
    sigma_sq = np.mean(epsilon**2 * D_res**2) / np.mean(D_res**2)**2
    se = np.sqrt(sigma_sq / len(Y))

    return theta_hat, se
```


Live Coding: Implementing Double ML

Exercise

Goal: Implement DML from scratch with cross-fitting.

Steps:

1. Implement cross-fitting loop
2. Use FLAML AutoML to find best models
3. Compute residuals for each fold
4. Run final line-fit on residuals
5. Compute standard errors
6. Compare with true elasticity!



Open notebook and follow along!

Key insight: DML recovers the true effect even when $f(X)$ is nonlinear!

The Overlap Analogue in Partial Linearity

Key Insight

The overlap condition becomes $\text{Var}(\tilde{D}) > \epsilon$

- ▶ $\tilde{D} = D - \mathbb{E}[D|X]$ is the unpredictable part of treatment
- ▶ We need this to have positive variance **on average**
- ▶ Much weaker than pointwise overlap!

For binary treatments: $\text{Var}(\tilde{D}) = \mathbb{E}[p(X)(1 - p(X))]$

- ▶ We need propensity bounded away from 0 and 1 **on average**
- ▶ Can extrapolate across covariate regions under partial linearity

Part VI: Interpretation with Mis-Specification

What Does DML Estimate if CEF is not Partially Linear?

Treatment Effect Heterogeneity: Binary Case

What if the true model has heterogeneous effects?

$$\mathbb{E}[Y|D, X] = \tau(X) \cdot D + f(X)$$

where $\tau(X)$ varies with X . We can still run the DML procedure. Irrespective of partial linearity, DML will converge to the quantity:

$$\theta_{\text{DML}} = \frac{\mathbb{E}[\tilde{D}\tilde{Y}]}{\text{Var}(\tilde{D})}, \quad \text{with } \tilde{Y} = Y - \mathbb{E}[Y | X] \text{ and } \tilde{D} = D - \mathbb{E}[D | X]$$

Overlap-Weighted Average Treatment Effect

DML estimates a weighted average of $\tau(X)$, with weights $\propto \text{Var}(D|X)$ —more weight where treatment is more “random” (less predictable), i.e.,

$$\theta_{\text{DML}} = \frac{\mathbb{E}[\tilde{D}\tilde{Y}]}{\text{Var}(\tilde{D})} = \frac{\mathbb{E}[\tau(X)\text{Var}(D|X)]}{\mathbb{E}[\text{Var}(D|X)]}$$

Why Use DML Even for Binary Treatments?

Even though DR works for binary treatments, DML has advantages:

1. **Lower variance:** Weights by $\text{Var}(D|X) = p(X)(1 - p(X))$, downweighting extreme propensities
2. **Natural interpretation:** Overlap-weighted ATE
3. **Robustness to overlap violations:** Works even with some regions of poor overlap

Trade-off

Has lower variance and does not require overlap, but requires partial linearity assumption (constant treatment effect). If effects are heterogeneous, DML gives a weighted average, not the population ATE.

Python Code for Binary Treatment DML

```
from sklearn.linear_model import LogisticRegression

def double_ml_binary(Y, D, X, n_folds=5):
    kf = KFold(n_splits=n_folds, shuffle=True)
    Y_res, D_res = np.zeros(len(Y)), np.zeros(len(Y))

    for train, test in kf.split(X):
        # Outcome model
        h_model = AutoML().fit(X[train], Y[train])
        Y_res[test] = Y[test] - h_model.predict(X[test])

        # Propensity model (classifier for binary D)
        p_model = AutoMLClassifier().fit(X[train], D[train])
        D_res[test] = D[test] - p_model.predict_proba(X[test])[:,1]

    theta_hat = np.mean(Y_res * D_res) / np.mean(D_res**2)
    return theta_hat
```

Continuous Treatments: Weighted Average Derivatives

For continuous treatments **without** partial linearity, let $g(d, X) = \mathbb{E}[Y|D = d, X]$. Even without assuming $g(d, X) = \theta d + f(X)$, DML estimates a **weighted average marginal effect**:

$$\theta_{\text{DML}} = \mathbb{E} \left[w(D, X) \cdot \frac{\partial}{\partial d} g(d, X) \Big|_{d=D} \right]$$

for some weight function $w(D, X)$.

Think of it as: “On average, how much does outcome change when we nudge treatment?”

The weights are complex in general, but DML still has a causal interpretation as a weighted average derivative!

Special Case: Gaussian Treatment

When $D|X \sim N(p(X), \sigma^2(X))$, the story simplifies beautifully. The DML numerator is:

$$\mathbb{E}[\tilde{Y} \cdot \tilde{D}] = \mathbb{E}[Y \cdot \tilde{D}] = \mathbb{E}[g(D, X) \cdot (D - p(X))]$$

By **Stein's Lemma** (a result for Gaussian random variables):

$$\mathbb{E}[g(D, X) \cdot (D - p(X))] = \mathbb{E}\left[\sigma^2(X) \cdot \frac{\partial}{\partial d} g(d, X) \Big|_{d=D}\right]$$

Variance-Weighted Average Marginal Effect

$$\theta_{\text{DML}} = \frac{\mathbb{E}[\sigma^2(X) \cdot \partial_d g(D, X)]}{\mathbb{E}[\sigma^2(X)]}$$

Interpretation of the Gaussian Case

Key Result

Under Gaussian treatment, DML estimates the **variance-weighted average marginal effect**:

$$\theta_{\text{DML}} = \frac{\mathbb{E} [\sigma^2(X) \cdot \partial_d g(D, X)]}{\mathbb{E} [\sigma^2(X)]}$$

Intuition:

- ▶ $\partial_d g(D, X)$ is the marginal effect at (D, X)
- ▶ $\sigma^2(X)$ is the conditional variance of treatment given X
- ▶ We weight by variance: more weight where treatment is more “experimental”

Even without partial linearity, DML has a clean causal interpretation!

Part VII: Industry Applications

Double ML in Practice

DML at Amazon

Amazon paper claims usage of DML in their experimentation platform for causal inference at scale.

Paper: “Double Machine Learning at Scale to Predict Causal Impact of Customer Actions” (Amazon Science)

Key applications:

- ▶ Estimating effect of customer actions on purchases
- ▶ Price elasticity estimation
- ▶ Marketing attribution

Why DML? Handles high-dimensional covariates, flexible ML models, valid inference, continuous treatments, lower variance than DR, less sensitive to overlap!

Software Implementations

EconML (Microsoft):

```
from econml.dml import LinearDML

est = LinearDML(model_y=AutoML(), model_t=AutoML())
est.fit(Y, D, X=None, W=X)
est.summary()
```

DoubleML (Python package):

```
from doubleml import DoubleMLPLR
from sklearn.ensemble import RandomForestRegressor

ml_g = RandomForestRegressor()
ml_m = RandomForestRegressor()
dml = DoubleMLPLR(data, ml_g, ml_m)
dml.fit()
dml.summary
```


Live Coding: Comparing with Libraries

Exercise

Goal: Compare our implementation with EconML and DoubleML.

Steps:

1. Run EconML LinearDML
2. Run DoubleML DoubleMLPLR
3. Compare all estimates:
 - ▶ OLS (linear controls)
 - ▶ Our DML implementation
 - ▶ EconML
 - ▶ DoubleML
4. Interpret as price elasticity!



Open notebook and follow along!

Summary: Four Key Takeaways

1. **The Frisch-Waugh-Lovell Theorem:** The coefficient θ in $\text{OLS}(Y \sim D, X)$ equals the slope from regressing the “surprise” in Y on the “surprise” in D —the normalized correlation of residuals.
2. **Insensitivity to First-Stage Errors:** When computing standard errors for θ , we can ignore estimation error in the first-stage regressions—another “insensitive formula” like the DR estimator.
3. **Partial Linearity Enables Flexibility:** Assuming $\mathbb{E}[Y|D, X] = \theta D + f(X)$ allows ML for $f(X)$ while still getting valid inference on θ .
4. **Double ML and Weighted Average Effects:** DML estimates variance-weighted average treatment effects (or marginal effects under Gaussian treatments via Stein’s lemma)—widely used in industry.

Next Time

- ▶ Beyond partial linearity?
- ▶ General principle and framework of “insensitivity” for causal ML estimation problems

References

- ▶ Chernozhukov, V., et al. (2018). Double/debiased machine learning for treatment and structural parameters. *The Econometrics Journal*, 21(1), C1-C68.
- ▶ Frisch, R., & Waugh, F. V. (1933). Partial time regressions as compared with individual trends. *Econometrica*, 1(4), 387-401.
- ▶ Lovell, M. C. (1963). Seasonal Adjustment of Economic Time Series and Multiple Regression Analysis. *Journal of the American Statistical Association*, 58(304), 993-1010.
- ▶ Amazon Science. Double machine learning at scale to predict causal impact of customer actions.
- ▶ EconML Documentation: <https://econml.azurewebsites.net/>
- ▶ DoubleML Documentation: <https://docs.doubleml.org/>